



# MuseumSpark: React + ChatGPT-Powered Trip Planner on Azure

**MuseumSpark** is a personal web application for planning museum trips, combining a React frontend with a Python backend that integrates OpenAI's ChatGPT for dynamic content. The stack uses a modern toolchain: **React** (bootstrapped with Vite) for the web app UI, **Tailwind CSS v4** for design, a **FastAPI**-based Python API (leveraging Pydantic AI agents for ChatGPT calls), and **SQLite** as a lightweight database. All components are self-hosted on an Azure Windows Server VM under your control. This architecture provides an end-to-end solution for interactive trip planning – from a responsive user interface to on-demand AI-generated itineraries – while remaining simple enough for a hobby project.

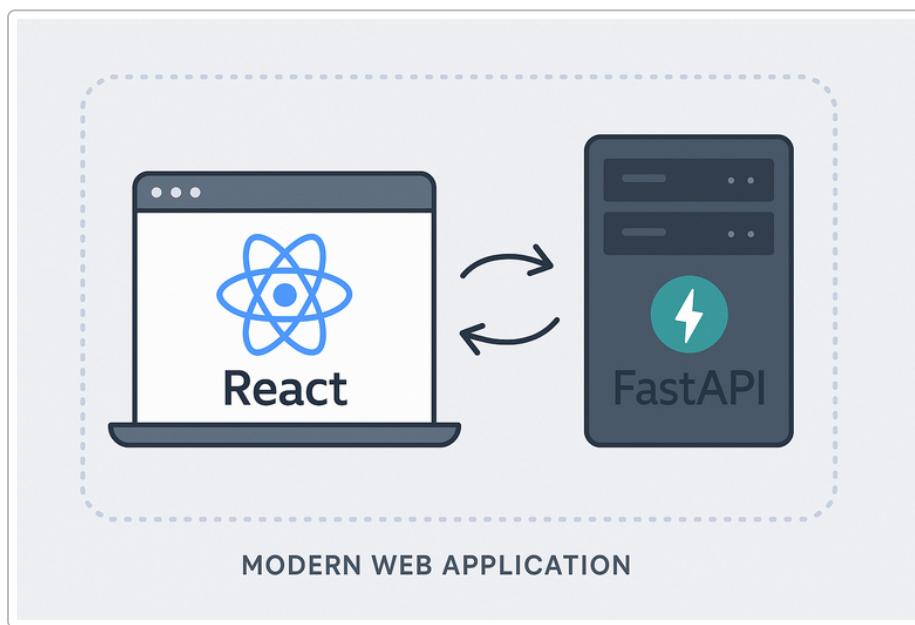


Figure: High-level architecture of MuseumSpark – a React single-page app (SPA) communicates with a FastAPI (Python) backend. The backend handles user requests, interacts with the SQLite database, and makes on-demand calls to the ChatGPT API. Both the frontend and backend are deployed on an Azure Windows VM, simplifying infrastructure and avoiding CORS issues by serving the SPA and API from the same origin.

## Tech Stack Overview

- **React + Vite:** The frontend is built with React, using Vite as the build tool. Vite provides a “blazing fast frontend build tool” experience <sup>1</sup> with instant hot-reloads and optimized production bundling. This ensures a smooth developer workflow and efficient packaging of the MuseumSpark web app.
- **Tailwind CSS v4:** Tailwind is used for styling the UI with utility-first CSS classes. Tailwind v4 integrates seamlessly with Vite and React, allowing rapid design iteration without writing custom CSS. The utility classes keep the site’s design consistent and easily maintainable.

- **Python FastAPI Backend:** The backend is a Python web API built with FastAPI. FastAPI is chosen for its modern design (built on Pydantic and type hints) and high performance, which makes it ideal for building APIs quickly <sup>2</sup>. It will expose endpoints for user authentication, trip management, and triggering ChatGPT queries.
- **Pydantic AI (ChatGPT Agent):** MuseumSpark's standout feature is AI-assisted planning. We use **Pydantic AI** to manage ChatGPT integration. Pydantic AI is a framework by the Pydantic team that simplifies building “*production-grade AI agents*” and provides type-safe interaction with LLMs <sup>3</sup> <sup>4</sup>. This means we can define an *Agent* that calls OpenAI’s GPT model (e.g., GPT-4) and incorporate the responses directly into our app logic. Pydantic AI supports many LLM providers (including OpenAI, Anthropic, etc.) <sup>4</sup>, ensuring flexibility.
- **SQLite Database:** We use a local SQLite database to store persistent data like user accounts and saved trips. SQLite is file-based and “*perfect for development*” or low-traffic apps <sup>5</sup> – no server setup required. It strikes a good balance between simplicity and functionality for this hobby project, given that concurrency demands are low.
- **Azure Windows VM Deployment:** The entire application (frontend and backend) is hosted on a Windows Server Virtual Machine on Azure, which you manage. This gives full control over the environment. The React app will be served as static files (either via the Python server or IIS), and the FastAPI app will run (via Uvicorn or similar) to handle API requests. Running on a single VM keeps deployment simple: you only need to manage one host for both UI and API.

With this stack, MuseumSpark remains lightweight yet powerful – well-suited for a personal project that isn’t aiming for commercial scale. Next, we’ll dive into how to set up each piece of the stack and get everything working together under the **MuseumSpark** brand.

## Frontend Setup: React, Vite, and Tailwind CSS

Setting up the MuseumSpark frontend involves creating a React application with Vite and integrating Tailwind CSS for styling. Below are the key steps and considerations:

1. **Initialize the Vite + React project:** Use Vite’s starter to scaffold a new React app. In your terminal, run:

```
npm create vite@latest museumspark-web -- --template react
```

This creates a new React project (in the `museumspark-web` directory) configured with Vite’s build system. Vite’s dev server will enable lightning-fast hot module replacement and a smooth development experience <sup>6</sup>.

2. **Install Tailwind CSS:** Change into the project directory and install Tailwind and its Vite plugin:

```
cd museumspark-web
npm install tailwindcss @tailwindcss/vite postcss autoprefixer -D
```

After installation, initialize Tailwind config by running `npx tailwindcss init` (which creates `tailwind.config.js`). You’ll also want to create a `postcss.config.js` with the Tailwind and

autoprefixer plugins. This setup allows Tailwind to scan your files and generate the necessary CSS. (*Tailwind's official Vite guide covers adding the plugin and config files* <sup>7</sup> <sup>8</sup>.)

3. **Configure Vite for Tailwind:** Update the Vite config (`vite.config.js` or `.ts`) to use the Tailwind plugin. For example:

```
// vite.config.js
import { defineConfig } from 'vite';
import react from '@vitejs/plugin-react';
import tailwindcss from '@tailwindcss/vite';
export default defineConfig({
  plugins: [react(), tailwindcss()],
});
```

This ensures Vite processes your CSS with Tailwind. In your main CSS (e.g., `index.css` or `App.css`), import Tailwind's base styles:

```
@tailwind base;
@tailwind components;
@tailwind utilities;
```

Now you can use Tailwind classes in your JSX. For example, a heading can have `className="text-3xl font-bold underline"` and Tailwind will apply the appropriate styles once built <sup>9</sup>.

4. **Developing and Building:** During development, run `npm run dev` to start Vite's dev server <sup>10</sup>. You can view the app at `http://localhost:5173` (default Vite port) and enjoy instant feedback as you modify your React components. Once you are happy with the app, generate a production build with:

```
npm run build
```

This outputs static files (HTML, JS, CSS, assets) to a `dist/` folder by default. These files will be deployed to the server. Vite optimizes and minifies the bundle for performance, producing an efficient build ready for deployment.

5. **Using Tailwind in MuseumSpark:** Tailwind v4 brings utility classes for rapid UI building. For MuseumSpark, this means you can easily create a clean, responsive layout for trip listings, forms, and buttons without writing custom CSS. Tailwind's approach scans your React components for class names and "*generates the corresponding styles then writes them to a static CSS file*" at build time <sup>11</sup>, so the production app only includes the CSS actually used (keeping it lightweight). For example, using classes like `grid md:grid-cols-2 gap-4` can quickly create a responsive gallery of museum cards.

In summary, the frontend setup gives MuseumSpark a modern, branded interface. The combination of React's component-based UI, Vite's fast build system, and Tailwind's utility classes means development is efficient and the end result is a snappy app that looks polished. Next, we connect this frontend to a powerful backend that can handle user data and AI-generated content.

## Backend API Design: FastAPI and Pydantic AI Integration

The backend of MuseumSpark is a Python API server responsible for all the heavy lifting: serving data from the database, authenticating users, and integrating with the ChatGPT API via a Pydantic AI agent. We use **FastAPI** to structure the web API, and Pydantic AI to simplify calls to the OpenAI GPT model. Here's how to set it up:

- **FastAPI Framework:** FastAPI is ideal for this project due to its speed and developer-friendly design. It uses Python type hints and Pydantic under the hood to validate request and response data. This means we can define clear data models (for example, a `Trip` or `User` schema) and FastAPI will automatically parse incoming JSON into Python objects and enforce types. FastAPI also automatically generates interactive docs (Swagger UI), which is handy even in a personal project to test API endpoints. To start, install FastAPI and Uvicorn (an ASGI server) in your Python environment:

```
pip install fastapi uvicorn pydantic-ai
```

(We include `pydantic-ai` here as well, more on that next.)

- **Defining API Endpoints:** We will have endpoints for things like user registration & login, creating or retrieving trips, adding a new city to a trip plan, and refining a trip's details. For example:

- `POST /auth/register` and `POST /auth/login` – to create an account or authenticate (returning a token).
- `GET /trips` and `POST /trips` – to fetch all saved trips for the logged-in user, or create a new trip.
- `POST /trips/{id}/add_city` – to add a new city/destination to an existing trip (this could trigger a ChatGPT call to fetch info about that city's museums).
- `POST /trips/{id}/refine` – to refine or update the trip plan (likely also involving an AI call).
- `GET /trips/{id}` – to retrieve details of a specific trip (possibly including any AI-generated itinerary details stored).

Each endpoint will use FastAPI's dependency injection to ensure the user is authenticated (via a token or session) and to provide database access. FastAPI makes it easy to declare dependencies like a `get_current_user` function (secured by OAuth2 token) that runs before protected endpoints. This keeps our logic clean.

- **Using Pydantic AI for ChatGPT:** Pydantic AI allows us to create an **Agent** that interfaces with OpenAI's API in a structured way. At startup, we can initialize an agent with the desired model, e.g.:

```
from pydantic_ai import Agent
agent = Agent("openai:gpt-4") # or gpt-3.5-turbo, etc.
```

This agent object will handle sending prompts to the OpenAI API and returning responses. Pydantic AI is “*a Python framework that acts as a bridge between developers and LLMs*”, letting us define tools, system prompts, and expected output schemas for the AI <sup>3</sup> <sup>12</sup>. For MuseumSpark, we might set a system prompt for the agent like: “*You are MuseumSpark, an AI assistant that provides travel itineraries focusing on museums.*” When a user adds a city to their trip, our endpoint can call something like:

```
prompt = f"Give a short travel itinerary for museums in {city_name} for a
first-time visitor."
result = agent.run_sync(prompt)
itinerary_text = result.data # the AI's response
```

We can then send this `itinerary_text` back to the frontend or store it in SQLite for later.

*Why Pydantic AI?* – It provides nice features like type validation of AI outputs (we can define a Pydantic model for the expected response format), and easy handling of conversation context. In fact, the Pydantic AI FastAPI chat example demonstrates maintaining chat history across requests to give the model context <sup>13</sup>. We can leverage similar ideas if we allow iterative refinement: e.g., store the last AI response and user feedback so that `agent` can incorporate it in the next prompt to refine the itinerary.

- **FastAPI and Pydantic Integration:** Because FastAPI and Pydantic (and by extension Pydantic AI) are designed to work together, adding the agent to our app is straightforward. We might create a dependency that provides the `agent` to any path operation that needs it, or simply use it within route functions. For example:

```
@app.post("/trips/{id}/add_city")
async def add_city(id: int, city: CityInput, user: User =
Depends(get_current_user)):
    # ... (code to add city to DB)
    # Call AI for itinerary snippet
    ai_prompt = f"Suggest 3 must-see museums in
{city.name} with a one-line description each."
    result = agent.run_sync(ai_prompt)
    suggestions = result.data
    # Save suggestions to DB, or attach to trip object
    return {"city": city.name, "suggestions": suggestions}
```

In this way, the heavy operation of fetching new data is offloaded to ChatGPT on-demand. The Pydantic AI agent handles communication with the OpenAI API under the hood, so our code stays clean and synchronous-looking (using `run_sync` for simplicity in this example).

- **OpenAI API Key Management:** You will need an OpenAI API key to use ChatGPT. Make sure to keep this secret. On the Azure VM, you can set an environment variable (e.g., `OPENAI_API_KEY`) so that the Pydantic agent or the OpenAI library can read it. Pydantic AI by default may use OpenAI's environment variables or you can configure the agent with an API key in code (depending on the version). It's important not to expose this key on the client side – which we avoid by doing all AI calls on the server.

In summary, the backend is built to handle all **MuseumSpark** operations via a defined API. FastAPI gives us a robust framework for the HTTP layer and data handling, while the Pydantic AI agent seamlessly integrates ChatGPT's capabilities. The result is an API that can create personalized museum itineraries on the fly and serve them to the React frontend.

## Database and User Accounts with SQLite

For persistence, MuseumSpark uses a SQLite database. This choice aligns with the project's lightweight, non-commercial nature by keeping things simple and self-contained:

- **SQLite Setup:** SQLite operates with a single file (e.g., `museumspark.db`) on disk. Python's standard library includes SQLite support (`sqlite3` module), but you might find it easier to use an ORM like **SQLAlchemy** for defining models and interacting with the database. For instance, using SQLAlchemy you can define a connection string `DATABASE_URL = "sqlite:///museumspark.db"` and create an Engine with `connect_args={"check_same_thread": False}`<sup>14</sup>. The `check_same_thread=False` flag is important because it tells SQLite to allow access from multiple threads – FastAPI's server (Uvicorn) might handle requests in different threads, and without this SQLite defaults to single-threaded mode<sup>15</sup>. After engine setup, create a session factory and define your models (as classes inheriting from `Base`). On app startup, you can call `Base.metadata.create_all(engine)` to create tables if they don't exist.
- **Data Models:** We'll have at minimum two tables in the database – one for **users** and one for **trips** (plus possibly a table for trip destinations or itinerary items). For example, a `User` model might have fields: `id` (Integer, primary key), `username` (string, unique), and `password` (string, hashed). A `Trip` model could have `id`, `user_id` (owner, foreign key to User), `title` (name of the trip plan), etc. If we store AI-generated itinerary data, that could be another table or a JSON field in `Trip` or a related `ItineraryItem` table.
- **Lightweight User Accounts:** Although this is a hobby project, we still want basic authentication so each user can save their own trips. "Lightweight" in this context means we'll avoid overly complex setups (like social logins or multi-factor auth) and stick to a simple username/password system with token auth. We can implement this as follows:
  - **User Registration:** An endpoint `POST /auth/register` will accept a username and password. The password should be hashed before storage. We can use a library like `passlib[bcrypt]` for hashing. *Never store plain passwords!* FastAPI tutorials emphasize storing only hashed passwords<sup>16</sup> for security. When a new user registers, hash the password (e.g., using bcrypt) and save the user record in SQLite.

- **User Login:** An endpoint `POST /auth/login` can use FastAPI's OAuth2PasswordRequestForm (which expects `username` and `password` fields) to parse credentials. We then verify the password against the stored hash (using passlib's verify function) <sup>17</sup>. If valid, we create a JWT (JSON Web Token) or an equivalent token to return to the client. FastAPI's `OAuth2PasswordBearer` can be used to parse and check this token on subsequent requests. The token contains the user's identity and possibly an expiry. For simplicity, you could use a short JWT validity (e.g., 1 day) and a secret key stored as an environment variable on the server. The Medium guide on FastAPI + JWT with SQLite provides a full example of this setup <sup>18</sup> <sup>19</sup>.

- **Auth on API Calls:** Once the client (React app) gets the token after login, it will include it in the Authorization header (`Bearer <token>`) for future API calls. FastAPI will use a dependency (`OAuth2PasswordBearer`) to fetch the token and a function to decode and validate it. If valid, the corresponding user is loaded (e.g., query the SQLite DB for that user ID) and attached to the request context. We can then protect routes like `/trips` by requiring an authenticated user. Given the low user count expected, this simple JWT approach is sufficient and much lighter than setting up a full OAuth server.

- **Saving Trips and Data:** The **Trips** data that users create (destinations, itinerary suggestions, etc.) will be stored in SQLite as well. For example, when a user adds a new trip, a row is inserted into `trips` with the title and `user_id`. If they add a city to a trip and we fetch AI suggestions, we can store those suggestions in a separate table (e.g., `attractions` table with `trip_id`, city name, description). Alternatively, since this is a personal app, storing a blob of text or JSON with the itinerary in the trip record could be acceptable. SQLite can even store JSON natively in a TEXT field (or using SQLite's JSON extension) if you want to keep the AI response structured.

- **Concurrency Considerations:** For a personal site, it's likely only one user (or a handful) will use it at a time, so SQLite will handle the load easily. However, remember that SQLite allows only one write operation at a time. FastAPI might handle requests concurrently, so it's wise to use a single database connection per request (SessionLocal pattern with SQLAlchemy, or opening a new sqlite3 connection per request if using raw sqlite3). This avoids conflicts. The app should also handle errors like `database is locked` gracefully (e.g., by retrying or serializing certain operations) if they ever occur under concurrent use.

Overall, using SQLite and simple user accounts aligns with MuseumSpark's needs: minimal setup, easy to maintain, and enough to provide personalization. The database provides a place to keep user preferences and past AI-generated results, so that each user's experience is saved between sessions.

## Dynamic Trip Planning with On-Demand ChatGPT Calls

One of the defining features of MuseumSpark is its ability to generate trip itineraries and suggestions on the fly using ChatGPT. Rather than pre-populating a database with museum data, the app makes **on-demand API calls** to fetch or refine information. Here's how this dynamic content generation works and how we ensure it supports all required functionality:

- **On-Demand API Calls:** When a user interacts with the MuseumSpark app – for example, by adding a new city to their trip – the frontend triggers an API call to the backend (e.g., `POST /trips/{id}/add_city` as discussed). The backend, upon receiving the request, invokes the ChatGPT agent to

get relevant information about that city's museums. This design means **we only fetch data when the user needs it**, keeping the app's initial data load small. It's a pull-based model: data is pulled in from the AI as required, rather than pushing a whole database of content upfront.

- **Use Cases:** The AI calls can support various functionality:
  - **Adding a New City:** The user provides a city name, and the backend asks ChatGPT something like, *"List the top 5 museums in [City] and a brief note about each."* The result is a curated list of attractions which is then shown to the user as suggestions. The user might then select which ones to include in their trip plan.
  - **Refining Data:** Suppose the user wants a more detailed plan (e.g., opening hours or a 2-day itinerary). When they click "Refine itinerary" (or similar), the frontend might call `POST /trips/{id}/refine` with parameters (like "2-day detailed plan"). The backend could send the current plan and user request to ChatGPT: *"Using the current trip plan as a base, create a 2-day schedule covering all listed museums, with optimal visiting order."* The AI's answer could then replace or augment the stored itinerary. This iterative refinement is possible because we maintain context: the backend can include previous AI outputs or user notes in the prompt.
  - **Answering Questions:** If implemented, users could ask free-form questions (e.g., "What's a famous art piece to see in the Louvre?"). The backend can relay this question to ChatGPT and return the answer. This turns MuseumSpark into an interactive travel assistant.
  - **Pydantic AI Agent Advantages:** By using the Pydantic AI agent, we gain some structure around these calls. We could define, for instance, a Pydantic model for the output of certain prompts (e.g., a model with `museums: List[MuseumRecommendation]` where `MuseumRecommendation` has fields like `name` and `description`). The agent can be configured with this `result_type`, so it will attempt to format ChatGPT's answer into that shape. This way, if the AI returns something unexpected, the Pydantic validation might catch it (marking an error or `UnexpectedModelErrorBehavior`), which we can handle. While not foolproof, it adds reliability for production-grade usage <sup>3 12</sup> – useful even in a hobby project to avoid crashing on odd AI outputs.
  - **Maintaining Context:** To support refinement and follow-up questions, you'd want to maintain conversation context with ChatGPT. Pydantic AI can manage chat histories (as shown by their FastAPI chat example) by storing previous messages <sup>13</sup>. In practice, you might keep the last prompt and response in the database (tied to the trip or user session) and then, when a refine request comes in, feed them as part of the new prompt (as system or assistant messages). This helps ChatGPT remember what it suggested before, leading to more coherent refinements. However, be mindful of token limits – you wouldn't send an entire conversation history if it's not needed.
  - **Rate Limits and Performance:** Since every new city or refinement triggers an API call to OpenAI, consider the rate limits and latency. For a personal project with few users, hitting OpenAI's rate limits is unlikely, but it's good to handle exceptions (e.g., API returns 429 or errors). You might implement a simple cache: if the same user asks for "museums in Paris" twice, you can store the result from the first call in SQLite and return it on subsequent requests without calling the API again. This saves time and cost. Similarly, if refining, you might only call the API if the request or context actually changed.

- **Cost Management:** Using the API on demand is cost-efficient: you're not charged unless the user actually requests something. Still, it's wise to use a model appropriate for the task – maybe `gpt-3.5-turbo` for quick list generation (since it's cheaper and fast), and only use `gpt-4` if you need more detailed or higher quality outputs. Pydantic AI allows specifying different models or even providers, so you have flexibility to switch as needed <sup>4</sup>. For example, you might configure the agent to use GPT-3.5 by default and have a parameter for using GPT-4 on demand.

In essence, **MuseumSpark's AI integration** ensures that the content is fresh and tailored to the user's requests. By making ChatGPT calls on demand, the app remains lean (no heavy initial data load) yet rich in information. The backend design with Pydantic AI provides a structured way to harness this power, which keeps our implementation maintainable.

## Deployment on Azure Windows VM

Deploying both the frontend and backend on an Azure Windows virtual machine involves setting up the environment, serving the static files for the React app, and running the FastAPI server continuously. Here's a step-by-step guide to get MuseumSpark up and running on your Azure VM:

1. **Provision the Azure VM:** You likely already have the Windows Server VM available. Ensure it has sufficient resources (MuseumSpark is lightweight – even a small VM with 1-2 CPU cores and 2-4 GB RAM should be plenty for a personal app). Make sure you have administrator access to the VM via Remote Desktop. Also, check networking settings: you'll want to allow inbound traffic on the port you plan to serve your app (port 80 for HTTP is standard, or you can use another port like 8000). In Azure, this means configuring the Network Security Group for the VM to permit that port, and likewise opening the port in the Windows Firewall.
2. **Install Required Software:** On the VM, install:
  3. **Node.js** (if you plan to build or re-build the frontend on the server). Download the LTS Windows installer from [nodejs.org](https://nodejs.org) and run it. This will give you Node and npm.
  4. **Python 3.x** (if not already installed on the server). Use the latest stable release (e.g., Python 3.11+). During installation, check "Add Python to PATH" for convenience.
  5. Optionally, **Git** (if you want to pull your code from a repository). If you prefer, you can also use RDP's copy-paste or FTP to transfer files without git.
6. Alternatively, ensure you have your project files (frontend `dist` folder and backend code) ready to copy to the server.
7. **Deploy the Frontend (Static Files):** After building the React app (you can build on your local machine and copy the `dist` folder, or do `npm install && npm run build` on the VM), you need to serve these static files:
  8. *Option A: Serve via FastAPI:* This is a simple and recommended approach. We can configure FastAPI to serve the React build. Copy the entire `dist/` directory (containing `index.html` and static assets) to a known location on the server, say `C:\museumspark\frontend\dist`. In your FastAPI app code, add:

```
from fastapi.staticfiles import StaticFiles  
app.mount("/", StaticFiles(directory="frontend/dist", html=True),  
name="static")
```

This mounts the React app at the root URL. With `html=True`, any path not found will fallback to serving `index.html` <sup>20</sup> – critical for client-side routing to work. In effect, any request that isn't an API endpoint (like `/api/...`) will return the React app's page, allowing React Router to take over <sup>21</sup>. This way, you run a single server for both API and UI.

9. **Option B: Serve via IIS or another web server:** If you prefer using IIS (the built-in Windows web server), you could place the files in `C:\inetpub\wwwroot\` or set up a site pointing to the `dist` folder. However, you'd then also need to configure a reverse proxy for the API (so that calls to, say, `/api/...` on IIS get forwarded to the FastAPI app on a different port). This is more complex and usually unnecessary for a small app. Still, if you require using port 80 on IIS, one trick is to serve static files via IIS and run FastAPI on another port (like 8000), then use URL Rewrite or Application Request Routing in IIS to proxy API calls. For most cases, Option A is simpler.
10. **Deploy the Backend (FastAPI app):** Copy your backend code to the server (for example, `C:\museumspark\backend\` containing your Python files). On the server:

11. Create a Python virtual environment:

```
python -m venv C:\museumspark\venv
```

Activate it (for Windows Command Prompt: `C:\museumspark\venv\Scripts\activate.bat`; for PowerShell: `.\venv\Scripts\Activate.ps1`). Then install dependencies:

```
pip install fastapi uvicorn pydantic-ai sqlite3 sqlalchemy passlib[bcrypt]  
python-jose[cryptography]
```

(The latter packages are if you implement JWT auth as discussed. Adjust to your actual requirements.)

12. Ensure your environment variables are set. In Windows, you can set them via System Properties or in the command shell. At minimum set `OPENAI_API_KEY`. If using JWT, set a `SECRET_KEY` for signing.
13. **Run the app:** For testing, you can manually start Uvicorn. For example:

```
uvicorn main:app --host 0.0.0.0 --port 8000
```

(This assumes your FastAPI instance is named `app` in `main.py`; adjust module/name as needed.) If you want the app accessible on default HTTP port 80, you could use `--port 80` (but that might require admin privileges). Initially, try port 8000 to verify functionality. Go to `http://<your-vm-ip>:8000` in a browser – you should see the MuseumSpark React app load (served by FastAPI), and API endpoints like `http://<vm-ip>:8000/docs` should show the interactive docs if not disabled.

14. **Keep the app running:** Closing the RDP or terminal will stop Uvicorn if it's running in a console. For production, set up a persistent service. On Windows, one approach is to use **NSSM (Non-Sucking Service Manager)** to wrap the Uvicorn command in a Windows service. Install NSSM, then run something like:

```
nssm install MuseumSparkAPI "C:\museumspark\venv\Scripts\python.exe" "C:\museumspark\venv\Scripts\uvicorn.exe" "main:app" "--host" "0.0.0.0" "--port" "8000"
```

This creates a Windows service named "MuseumSparkAPI" that will run the Uvicorn server in the background. You can configure it to start automatically on boot and restart on failure. Alternatively, you could use a simpler approach like a scheduled task at startup or just remember to run it via RDP when needed (not ideal, but okay for a hobby project if occasional downtime is acceptable).

15. **Security:** Since this VM is exposed to the internet, ensure basic security. Use strong passwords for RDP. Limit the open ports (maybe only 80/8000 and 443 if you set up SSL). Speaking of SSL – for a personal project, you might not have a TLS certificate, but it's recommended if you use real user accounts. You can obtain a free Let's Encrypt certificate even on an Azure VM (though the automation on Windows is a bit trickier than Linux). At least, if you have a custom domain for MuseumSpark, you can configure it and consider securing it. If not, perhaps you'll run it as an HTTP-only site for just your own use or testing.
16. **Testing the Deployed App:** Once both the static files and API are being served on the VM, test the whole flow:

17. Navigate to `http://<your-vm-ip>:8000` (or whichever port/domain is set). The React app should load.
18. Try registering a new user via the UI, logging in, adding a trip, etc. All those actions will trigger API calls to the same host.
19. Verify that ChatGPT calls work. You might want to monitor the backend logs (Uvicorn console output) while testing. The first AI call might take a moment (due to model response time), but subsequent ones should function as expected.
20. If something isn't working (e.g., if API calls from the React app get blocked), it could be a CORS issue – but since we served React from the same origin, requests to the API on the same domain/port should be fine (no CORS preflight needed). CORS was more of a consideration during local development (React dev server to local API) and not in production when unified. FastAPI's docs on CORS and using CORSMiddleware are useful if needed <sup>22</sup>.
21. **Branding and Presentation:** Because the prompt says "Brand this MuseumSpark", ensure the deployment reflects the name. Set a proper title in the HTML (e.g., `<title>MuseumSpark</title>`). If using Swagger UI for the API docs, you can customize the FastAPI title and description to say "MuseumSpark API". Small touches like a favicon (perhaps a museum icon) served at `/static/favicon.ico` can make the app feel polished. Since it's a personal site, you have full freedom to style it to your liking with Tailwind – maybe a distinctive color scheme or logo that reflects "spark your museum journey"  (for fun, as the name suggests sparking interest in museums).

By following these deployment steps, you'll have MuseumSpark running live on your Azure VM, accessible to you (and any friends or beta testers you share it with). The React frontend and FastAPI backend live together on the VM, which simplifies development and avoids issues like cross-origin requests. You can iterate on features locally, then deploy updates by rebuilding the frontend and updating the backend code on the server.

---

**In conclusion,** MuseumSpark's architecture demonstrates how modern web technologies can be combined to create a powerful yet compact application. We used React and Tailwind to craft a responsive UI, FastAPI and SQLite to handle data and users, and Pydantic AI to seamlessly integrate ChatGPT's intelligence into the experience. The entire stack is hosted on a manageable Azure Windows VM, under the MuseumSpark brand. This setup is well-suited for a hobby project: it minimizes cost and complexity while maximizing the ability to experiment with cutting-edge features (like AI integration). With everything in place, MuseumSpark can indeed "spark" your museum adventures by providing personalized trip plans and insights at the click of a button, all within a single cohesive application.

#### Sources:

- FastAPI Documentation – *Serving Static Files* 23 24 (for mounting React build in FastAPI).
- David Muraya, *Serving a React Frontend with FastAPI* – handling client-side routing and static file mount 20 21 .
- Tailwind CSS Official Docs – *Using Tailwind with Vite* 7 8 .
- Vite Official Site – Vite is “*a blazing fast frontend build tool*” 1 .
- Pydantic AI Documentation – Pydantic AI is “*a Python agent framework*” for GenAI by the Pydantic team 3 , compatible with providers like OpenAI 4 .
- Pydantic AI FastAPI Chat Example – demonstrates maintaining chat history for context 13 .
- Gopinath V, *FastAPI JWT Auth with SQLite* (Medium) – on setting up SQLite with SQLAlchemy (check\_same\_thread) 14 15 and hashing passwords (never store plain text) 16 .

---

1 6 Vite | Next Generation Frontend Tooling

<https://vite.dev/>

2 Pydantic AI - Pydantic AI

<https://ai.pydantic.dev/>

3 4 12 Pydantic AI: Agent Framework. PydanticAI is a Python Agent Framework... | by Bhavik Jikadara | AI Agent Insider | Medium

<https://medium.com/ai-agent-insider/pydantic-ai-agent-framework-02b138e8db71>

5 14 15 16 17 18 19 FastAPI JWT Authentication with SQLite Database for Secure APIs | by Gopinath V | Medium

<https://medium.com/@gopinath.v2507/fastapi-jwt-authentication-with-sqlite-database-for-secure-apis-11c899efe6cd>

7 8 9 10 11 Installing Tailwind CSS with Vite - Tailwind CSS

<https://tailwindcss.com/docs/installation/using-vite>

13 Chat App with FastAPI - Pydantic AI

<https://ai.pydantic.dev/examples/chat-app/>

[20](#) [21](#) [22](#) Serving a React Frontend Application with FastAPI

<https://davidmuraya.com/blog/serving-a-react-frontend-application-with-fastapi/>

[23](#) [24](#) Static Files - FastAPI

<https://fastapi.tiangolo.com/tutorial/static-files/>